

# Chapter 7 - Ex2: Adult Dataset - Full

- Adult Dataset được cung cấp bởi UCI (University of California, Irvine) được sử dụng để phát triển mô hình dự đoán Predictive Model Development.
- Bộ dữ liệu adult.data và adult.test chứa 48.842 mẫu và có 14 attributes/features. Dữ liệu này được dùng để xây dựng model dự đoán và kiểm tra một mẫu có thu nhập >50K USD hay không. ### Attribute Information:
- age: continuous.
- workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.
- fnlwgt: continuous.
- education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.
- education-num: continuous.
- marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.
- occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.
- relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.
- race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.
- sex: Female, Male.
- capital-gain: continuous.
- capital-loss: continuous.
- hours-per-week: continuous.
- native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.
- Class: >50K, <=50K.

## Yêu cầu:

- Đọc dữ liệu adult.data, tiền xử lý dữ liệu.
- Xem xét tính cân bằng giữa hai loại mẫu. Trực quan hóa. Nhận xét.
- Nếu 2 loại mẫu này không cân bằng, hãy chọn một phương pháp cân bằng dữ liệu và thực hiện. Trực quan hóa kết quả.

```
In [1]: # Link tham khảo: https://towardsdatascience.com/under-sampling-a-performance-booster-on-imbalance
```

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: # Đọc dữ liệu, kiểm tra sơ bộ ban đầu, trực quan hóa, tiền xử lý dữ liệu
adult_train = pd.read_csv("adult/adult.data", header=None)
```

```
In [4]: adult_train.head()
```

```
Out[4]:
```

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<=50K
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=50K
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50K
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50K

```
In [5]: adult_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
0      32561 non-null int64
1      32561 non-null object
2      32561 non-null int64
3      32561 non-null object
4      32561 non-null int64
5      32561 non-null object
6      32561 non-null object
7      32561 non-null object
8      32561 non-null object
9      32561 non-null object
10     32561 non-null int64
11     32561 non-null int64
12     32561 non-null int64
13     32561 non-null object
14     32561 non-null object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
In [6]: adult_train.to_csv("adult_data.csv")
```


```
In [7]: # Không có dữ liệu null
```

```
In [8]: # Đếm theo loại: hiếm, phổ biến
occ = adult_train[14].value_counts()
occ
```

```
Out[8]:  <=50K      24720
        >50K       7841
        Name: 14, dtype: int64
```

```
In [9]: plt.bar(occ.index.values, occ.values)
```

```
Out[9]: <BarContainer object of 2 artists>
```



```
In [10]: # Chuyển dữ liệu phân loại thành dạng numeric dùng Label encoder và dummy encoder
```

```
In [11]: y_train = adult_train[14]
        X_train = adult_train.drop([14], axis=1)
```

```
In [12]: X_train.head(2)
```

```
Out[12]:
```

	0	1	2	3	4	5	6	7	8	9	10	11	12	13
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States

```
In [13]: y_train[:2]
```

```
Out[13]: 0      <=50K
        1      <=50K
        Name: 14, dtype: object
```

```
In [14]: from sklearn.preprocessing import LabelEncoder
```

```
In [15]: label_encoder = LabelEncoder()
        y_train_l = label_encoder.fit_transform(y_train)
```

```
In [16]: y_train_l[:2]
```

```
Out[16]: array([0, 0])
```

```
In [17]: # Categorical boolean mask
categorical_feature_mask = X_train.dtypes==object
# filter categorical columns using mask and turn it into a list
categorical_cols = X_train.columns[categorical_feature_mask].tolist()
categorical_cols
```

```
Out[17]: [1, 3, 5, 6, 7, 8, 9, 13]
```

```
In [18]: X_train_d = pd.get_dummies(data=X_train, columns=categorical_cols, drop_first=True)
```

```
In [19]: X_train_d.head(2)
```

```
Out[19]:
```

	0	2	4	10	11	12	1_ Federal- gov	1_ Local- gov	1_ Never- worked	1_ Private	...	13_ Portugal	13_ Puerto- Rico	13_ Scotland	13_ South	13_ Taiwan
0	39	77516	13	2174	0	40	0	0	0	0	...	0	0	0	0	0
1	50	83311	13	0	0	13	0	0	0	0	...	0	0	0	0	0

2 rows × 100 columns



## Áp dụng thuật toán với dữ liệu gốc

```
In [20]: from sklearn.linear_model import LogisticRegression
```

```
In [21]: model = LogisticRegression()
```

```
In [22]: model.fit(X_train_d, y_train_l)
```

```
Out[22]: LogisticRegression()
```

```
In [23]: from sklearn.metrics import confusion_matrix, accuracy_score
```

```
In [24]: y_pred = model.predict(X_train_d)
```

```
In [25]: accuracy_score(y_train_l, y_pred)
```

```
Out[25]: 0.7957679432449863
```

```
In [26]: cm = confusion_matrix(y_train_l, y_pred)
```

```
In [27]: cm
```

```
array([[23842, 878],
```

```
Out[27]:      [ 5772,  2069]], dtype=int64)
```

```
In [28]: # Đánh giá model  
from sklearn.metrics import classification_report, roc_auc_score, roc_curve
```

```
In [29]: print(classification_report(y_train_l, y_pred))
```

	precision	recall	f1-score	support
0	0.81	0.96	0.88	24720
1	0.70	0.26	0.38	7841
accuracy			0.80	32561
macro avg	0.75	0.61	0.63	32561
weighted avg	0.78	0.80	0.76	32561

```
In [30]: y_prob = model.predict_proba(X_train_d)  
y_prob
```


```
Out[30]: array([[0.45215974, 0.54784026],  
                [0.63260177, 0.36739823],  
                [0.80282842, 0.19717158],  
                ...,  
                [0.72915332, 0.27084668],  
                [0.78767922, 0.21232078],  
                [0.04969381, 0.95030619]])
```

```
In [31]: roc_auc_score(y_train_l, y_prob[:, 1])
```

```
Out[31]: 0.6291047901269116
```

```
In [32]: import matplotlib.pyplot as plt
```

```
In [33]: # calculate roc curve  
fpr, tpr, thresholds = roc_curve(y_train_l, y_prob[:, 1])  
# plot no skill  
plt.plot([0, 1], [0, 1], linestyle='--')  
plt.plot(fpr, tpr, marker='.')  
plt.show()
```



# Kết luận

- ROC\_AUC thấp
- precision class 1 cao nhưng recall thấp

# Undersampling

```
In [34]: from collections import Counter
         sorted(Counter(y_train_1).items())

Out[34]: [(0, 24720), (1, 7841)]

In [35]: # Vì Lượng dữ liệu class 1 tương đối nhiều => do đó ta sẽ áp dụng Undersampling
         # để giảm số mẫu của nhóm <=50k bằng với nhóm >50k

In [36]: from sklearn.utils import resample

In [37]: # có thể dùng cách resample

In [38]: data_train = X_train_d
         data_train[14] = y_train_1

In [39]: data_0 = data_train[data_train[14]==0]
         data_1 = data_train[data_train[14]==1]

In [40]: display(data_0.shape, data_1.shape)

(24720, 101)
(7841, 101)

In [41]: from sklearn.utils import resample

In [42]: data_0_resample = resample(data_0,
                                   replace = False, # sample without replacement
                                   n_samples = data_1.shape[0], # match minority n
                                   random_state = 27) # reproducible results

In [43]: downsampled = pd.concat([data_0_resample, data_1])
         downsampled.head()
```

Out[43]:

	0	2	4	10	11	12	1_ Federal-gov	1_ Local-gov	1_ Never-worked	1_ Private	...	13_ Puerto-Rico	13_ Scotland	13_ South	13_ Taiwan	13_ Thailand
31749	22	199426	10	0	0	17	0	0	0	1	...	0	0	0	0	
24093	31	91964	13	0	0	40	0	0	0	1	...	0	0	0	0	

	0	2	4	10	11	12	1_ Federal- gov	1_ Local- gov	1_ Never- worked	1_ Private	...	13_ Puerto- Rico	13_ Scotland	13_ South	13_ Taiwan	Thail
21539	37	60313	9	0	0	40	0	0	0	1	...	0	0	0	0	
24582	30	85708	9	0	0	40	0	0	0	1	...	0	0	0	0	
622	65	109351	5	0	0	24	0	0	0	1	...	0	0	0	0	

5 rows × 101 columns



```
In [44]: display(data_0_resample.shape, data_1.shape)
```

```
(7841, 101)
(7841, 101)
```

## Áp dụng thuật toán với dữ liệu Undersampling

```
In [45]: X_down = downsampled.drop(14, axis=1)
```

```
In [46]: y_down= downsampled[14]
```

```
In [47]: model_down = LogisticRegression()
```

```
In [48]: model_down.fit(X_down, y_down)
```

```
c:\program files\python36\lib\site-packages\sklearn\linear_model\_logistic.py:764: ConvergenceWarn
ing: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
```

```
Out[48]: LogisticRegression()
```

```
In [49]: y_pred_d = model_down.predict(X_down)
```

```
In [50]: accuracy_score(y_down, y_pred_d)
```

```
Out[50]: 0.6840326488968244
```

```
In [51]: cm = confusion_matrix(y_down, y_pred_d)
```

```
In [52]: cm
```

```
Out[52]: array([[5263, 2578],
```

```
[2377, 5464]], dtype=int64)
```

```
In [53]: print(classification_report(y_down, y_pred_d))
```

	precision	recall	f1-score	support
0	0.69	0.67	0.68	7841
1	0.68	0.70	0.69	7841
accuracy			0.68	15682
macro avg	0.68	0.68	0.68	15682
weighted avg	0.68	0.68	0.68	15682


```
In [54]: y_prob_d = model_down.predict_proba(X_down)
y_prob_d
```

```
Out[54]: array([[0.47990917, 0.52009083],
 [0.33516886, 0.66483114],
 [0.6269155 , 0.3730845 ],
 ...,
 [0.3878897 , 0.6121103 ],
 [0.56140787, 0.43859213],
 [0.03907586, 0.96092414]])
```

```
In [55]: roc_auc_score(y_down, y_prob_d[:, 1])
```

```
Out[55]: 0.7649199762119465
```

```
In [59]: fpr, tpr, thresholds = roc_curve(y_down, y_prob_d[:, 1])
# plot no skill
plt.plot([0, 1], [0, 1], linestyle='--')
plt.plot(fpr, tpr, marker='.')
plt.show()
```



## Kết luận

- ROC\_AUC cao hơn so với dữ liệu gốc
- precision class 1 gần như dữ liệu gốc và recall cao hơn #### => Áp dụng Undersampling dữ liệu cho kết quả tốt hơn so với dữ liệu gốc ban đầu.

```
In [ ]:
```



